

Identifying At-Risk Customers: A Predictive Analysis of Credit Card Attrition

1. Introduction

A. Overview

Customer attrition is still a major problem for companies in a variety of sectors, and its effects must be lessened with efficient tactics. Churn, also known as customer attrition, is the reduction in the number of customers who stop using a business's goods or services. This problem is widespread in sectors like financial services, e-commerce, and telecommunications that depend on subscription-based business models, where steady income streams are critical to the long-term viability of enterprises. Since keeping current customers is much less expensive than acquiring new ones, client retention is an essential part of successful business plans. As a result, anticipating and resolving churn has emerged as a crucial area of emphasis for customer relationship management and company performance.

The banking sector is no exception to these challenges. The growth of digital-only banks and technological breakthroughs have caused a dramatic change in the sector during the past 20 years. Customers today anticipate more customization, smooth digital integration, and measurable value from the services they use, which has significantly changed their expectations due to these changes. Customer churn is especially problematic in the banking industry because of its extensive effects. High churn rates affect the bank's capacity to establish and preserve enduring connections with clients in addition to causing immediate revenue losses.

One financial product where client attrition might have serious repercussions is credit card services. Credit card customers are often highly profitable for banks, contributing to revenue through transaction fees, interest payments, and auxiliary services. In addition to an immediate drop in revenue, banks often see a decline in customer lifetime value (CLV), a crucial indicator for long-term financial planning, when these clients depart.

Banks are increasingly using data-driven methods to comprehend and predict client loss in response to these difficulties. To address the unique requirements and preferences of at-risk clients, these actions could take the form of tailored offers, better customer support, or strengthened loyalty

programs. These goals may now be accomplished with the help of predictive analytics and machine learning models, which provide banks the ability to use massive datasets to find patterns and trends that guide retention tactics.

This study analyzes the BankChurners dataset to forecast credit card user attrition through the use of machine learning techniques. Numerous client factors, such as demographic information, account history, and transactional activities, are included in the collection. Through the analysis of this data, the study seeks to pinpoint the main causes of customer attrition and create prediction models that can precisely identify consumers who are at danger. Banks will be able to better manage resources and prioritize retention initiatives with the aid of these models.

This study aims to advance the general knowledge of attrition in the financial services industry by investigating the reasons and predictors of customer churn. It is anticipated that the results would emphasize the significance of contextual elements like customer tenure and demographic traits, as well as behavioral indicators like transaction frequency and credit utilization. The study lays the groundwork for focused interventions that tackle the root causes of attrition by identifying key churn drivers. Additionally, this analysis's application of machine learning techniques highlights the need of sophisticated analytics in addressing challenging business issues. The ability to anticipate and stop client attrition will continue to be a crucial factor in determining success as banks navigate an increasingly dynamic and competitive environment.

B. Literature review

In the case of credit card consumers, recent research has improved the accuracy of churn predictions by utilizing machine learning approaches. Al-Najjar et al. (2022) used five machine learning models—the Bayesian Network, C5 Decision Tree, CHAID Tree, CR Tree, and Neural Networks—as well as feature selection approaches to create a prediction model for credit card customer attrition. The study showed that combining multi-categorical factors enhanced model performance and underlined the significance of early prediction for retention measures. These

results are consistent with the emphasis on feature selection and sophisticated machine learning for churn prediction in this research.

Similarly, in order to determine when and why customers are likely to stop using their services, Siddiqui et al. (2023) created churn prediction models, concentrating on credit card users. Three models were suggested by the study: Model 1 included all variables, Model 2 distinguished between continuous and categorical features, and Model 3 used feature selection methods to rank the most important aspects. The study highlighted important characteristics of churners and determined the best machine learning techniques for churn prediction. In order to improve feature selection and churn prediction accuracy, the study also compared its models to other methods already in use. By applying several sampling techniques and building a customer churn prediction model with an enhanced XGBoost algorithm, Peng et al. (2023) addressed the problem of data imbalance in churn prediction. Interpretability analysis was also used in their study to offer practical advice to financial institutions looking to reduce customer churn.

All of these experiments demonstrate how well machine learning models predict credit card client attrition. As the main indications of churn, they emphasize the crucial relevance of behavioral measures like transaction frequency and credit utilization. Furthermore, improving the forecast accuracy and interpretability of these models depends critically on the use of feature selection strategies and the management of data imbalances.

Financial institutions can proactively identify clients who are at risk of leaving and put targeted retention plans into place by combining these cutting-edge analytical techniques. This contributes to the institution's long-term success by strengthening client relationships and reducing the possibility of revenue losses.

2. Dataset

A. Description and Exploratory Analysis

Key information about customer demographics, behavior, and account-related traits is revealed by the dataset analysis, which is crucial for comprehending attrition trends. With a mean age of almost 46 years, the age distribution shows that the majority of clients are between the ages of 40 and

60, indicating that the bank predominantly caters to middle-aged people. The gender breakdown indicates that female clients make up a slight majority. The "Graduate" and "High School" categories have the highest concentration of education levels. Also, according to income, the largest group makes less than \$40,000 per year, while data on marital status shows that about half of the clients are married.

Characteristics pertaining to behavior and accounts are essential to comprehending churn. Lower transaction volumes and counts increase the likelihood of customer attrition, underscoring the significance of engagement in retention. Customer who regularly use their accounts are less likely to depart, according to behavioral variables like total transaction count and total transaction amount, which have substantial negative associations with attrition. The average credit limit of attrited consumers is typically lower, which may restrict their financial options (Table 1).

Higher attrition rates are also linked to shorter tenures at the bank, as indicated by "Months on Book," as well. Short-term bank customers are less likely to stick around, which emphasizes the value of establishing lasting relationships. Additionally, another important metric linked to attrition is a lower "Total Relationship Count," which indicates the number of accounts a customer has with the bank. This highlights that clients are more likely to depart if they have fewer accounts or are less involved with the bank. Engagement is further shown as a critical predictor of attrition by declines in transactional activity, especially across consecutive quarters.

Furthermore, compared to current customers, attrited clients have higher average utilization percentages, which may be a sign of financial distress or discontent with their credit limitations. Low engagement is a key predictor of attrition, as seen by the clear patterns in the transactional behavior of attrited consumers. These clients are more likely to have low transaction counts, usually less than 20 transactions year, and low transaction quantities, with annual expenditure frequently dipping below \$2,000. Such little activity raises the risk of churn by indicating a decreased dependence on or level of satisfaction with the bank's services. Customers that have smaller open credit amounts or high utilization ratios are also more likely to experience attrition. These trends can be a sign of financial strain or discontent with the available

credit limits, which would encourage clients to look for other financial institutions.

Additional indications of disengagement among attrited clients can be found in changes in activity levels, especially between successive quarters. A notable decline in transaction counts and amounts from one quarter to the next (e.g., Q4 to Q1) underscores the role of diminishing engagement as a strong signal of potential churn. Predictive modeling is significantly hampered by the dataset's class imbalance. With only 16.1% of customers classified as "Attrited Customers" and 83.9% of customers classified as "Existing Customers," the dataset is significantly biased in favor of the majority class.

Mean	Attrited Customers	Existing Customers
Customer Age	46.66	46.26
Credit Limit	\$8,136	\$8,726
Transaction Amount	\$3,095	\$4,654
Transaction Count	44.93	68.67
Utilization Ratio	16.2%	29.6%

Table 1: Comparison of Key Metrics Between Attrited and Existing Customers

B. Preprocessing

The dataset was prepared for analysis in several ways. First, features that were not essential to the prediction task were eliminated, such as the unique customer identity CLIENTNUM. To ensure interoperability with machine learning algorithms, binary features for each category were created by transforming categorical variables—such as gender, marital status, and income category—using one-hot encoding. To standardize ranges and keep features with bigger scales from taking over the models, MinMaxScaler was used to normalize numerical variables like Total Transaction Amount and Average Utilization Ratio. To lessen their impact on model performance, outliers found during EDA were also capped using IQR criteria.

In order to improve the dataset, feature engineering was essential. Total Transaction Amount was divided by Months on Book to create a new variable called Transactions Per Month. Deeper insights into engagement patterns were made possible by this, which offered a normalized measure of client activity. The combination of these preprocessing steps ensured a robust and reliable dataset for modeling.

3. Predictive Task and Evaluation

The objective of the prediction task was to use the available features to categorize consumers as either existing or attrited. To give a thorough evaluation of model performance, suitable evaluation criteria were chosen in light of the dataset's notable class imbalance. To make sure that retention efforts are concentrated on truly at-risk consumers, precision was utilized to quantify the percentage of accurately identified attrited customers among all anticipated attrited customers. The model's capacity to identify at-risk instances was demonstrated by using recall to assess the percentage of correctly identified attrited consumers among all real attrited clients.

A balanced statistic to assess the model's overall performance was provided by the F1-score, which is a harmonic mean of precision and recall. Furthermore, the confusion matrix offers a thorough examination of true positives, false positives, true negatives, and false negatives, enabling error analysis and enhancing the model's dependability.

4. Model Development and Comparison

The development and comparison of models began with baseline approaches and advanced to more sophisticated techniques, including ensemble methods and feature selection, to improve performance. Each model's strengths and limitations were carefully assessed to achieve optimal predictions.

A. Baseline Models

Logistic Regression was used as the baseline model due to its simplicity and interpretability. This model provided a good starting point for comparison, as it is computationally efficient and easy to implement. However, its linear nature limited its ability to capture complex, non-linear relationships in the data. Logistic Regression achieved a precision of 0.78, recall of 0.48 for predicting attrition, and an F1-score of 0.69, demonstrating modest performance but highlighting the need for more advanced models.

K-Nearest Neighbors (KNN) offered slightly improved results compared to Logistic Regression in terms of precision of 0.75, but recall dropped to 0.4, and F1-score to 0.52, still leaving room for better solutions.

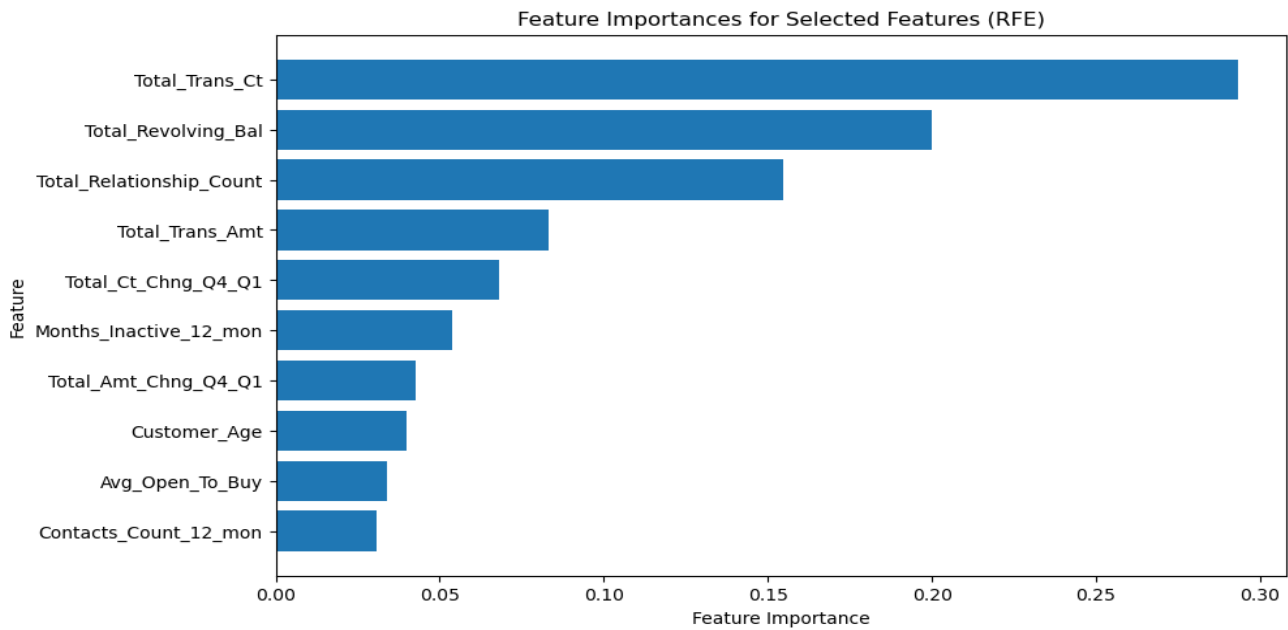


Fig 1.1 : Feature Importances for Selected Features Identified by Recursive Feature Elimination (RFE)

B. Advanced Models

Advanced machine learning models addressed the shortcomings of the baseline methods, and the performance significantly improved. With a precision of 0.80, recall of 0.81, and F1-score of 0.80, a Decision Tree model outperformed the baseline models and provided improved interpretability.

Random Forest (with SMOTE) tackled the class imbalance in the dataset by balancing the minority class using SMOTE before training. Random Forest achieved a precision of 0.87, recall of 0.90, and an F1-score of 0.89, demonstrating overall predictive power. Gradient Boosting further improved performance by reducing overfitting and capturing non-linear relationships more effectively than Random Forest. By focusing on errors from previous iterations, Gradient Boosting refined predictions achieving a precision of 0.93, a recall of 0.85, and an F1-score of 0.89 for the minority class, showcasing its ability to capture complex patterns in the data. LightGBM (LGBM), optimized for speed and scalability, offered an efficient alternative to Gradient Boosting. LGBM achieved a precision of 0.91, recall of 0.91, and an F1-score of 0.91, combining computational efficiency with strong predictive performance.

XGBoost emerged as the best-performing standalone model. We incorporated regularization techniques to prevent overfitting and underwent hyperparameter tuning through grid search to

optimize its performance. XGBoost achieved a precision of 0.92, a recall of 0.90, and an F1-score of 0.91, excelling across all evaluation metrics due to its robust handling of imbalanced data and complex relationships.

To streamline the modeling process and improve interpretability, Recursive Feature Elimination (RFE) was applied to XGBoost. This feature selection technique identified the most impactful predictors, with Total Transaction Count, Total Revolving Balance and Total Relationship Count emerging as the top features. By focusing on these critical variables, the model retained its strong predictive capabilities while reducing complexity. XGBoost with RFE achieved a precision of 0.92, recall of 0.91, and an F1-score of 0.91, balancing efficiency with performance.

The most important characteristics affecting customer attrition are highlighted by key findings from the investigation. Higher transaction counts were linked to a lower likelihood of turnover, making Total_Trans_Ct a powerful predictor of active involvement. While Avg_Utilization_Ratio demonstrated prudent credit usage and had a favorable correlation with retention, Total_Relationship_Count measured customer loyalty and retention. The importance of behavioral measures in churn prediction was highlighted by the fact that they routinely surpassed demographic indicators in predictive power, such as transaction and relationship count. These results were

Model	Precision (False/True)	Recall (False/True)	F1 Score (False/True)	Comments
Logistic Regression	0.78 / 0.93	0.62 / 0.97	0.69 / 0.95	Baseline model; struggled with recall for the False class.
KNN	0.73 / 0.89	0.40 / 0.97	0.52 / 0.93	Poor recall for the False class; not suitable for imbalanced datasets.
Decision Tree	0.80 / 0.96	0.81 / 0.96	0.80 / 0.96	Better balance than Logistic Regression and KNN.
Random Forest (SMOTE)	0.87 / 0.98	0.90 / 0.97	0.89 / 0.98	Handled class imbalance well using SMOTE.
Gradient Boost	0.93 / 0.97	0.85 / 0.99	0.89 / 0.98	Competitive performance but slightly below LightGBM and XGBoost models.
Light GBM	0.91 / 0.98	0.91 / 0.98	0.91 / 0.98	Excellent balance across metrics; tied with XGBoost in performance.
XG Boost	0.92 / 0.98	0.90 / 0.99	0.91 / 0.98	Robust performance and balanced metrics; highly interpretable.
XG Boost (RFE)	0.92 / 0.98	0.91 / 0.98	0.91 / 0.98	Best model; combines feature interpretability and top-tier performance.

Table 2 : Model Comparison Table: Precision, Recall, F1 Score, and Performance Analysis

confirmed by the use of Recursive Feature Elimination (RFE), which recognized these features as crucial and made it possible to create a simplified model with fewer variables while retaining high prediction accuracy.

5. Results and Conclusion

The best model for forecasting customer attrition, according to the data, is XGBoost with Recursive Feature Elimination (RFE). For the minority class (attrited consumers), the model's precision and recall were 0.92 and 0.91, respectively, indicating that it can effectively identify at-risk clients while reducing false positives. The model's dependability is further highlighted by the F1-score of 0.91, which shows a good balance between precision and recall. Furthermore, with precision, recall, and F1-scores all at 0.98, the model demonstrated outstanding performance for the majority class (current customers), guaranteeing no trade-off between the two classes. The model's resilience and consistency in managing unbalanced datasets are highlighted by its overall accuracy of 0.97.

References

1. AL-Najjar D, Al-Rousan N, AL-Najjar H. Machine Learning to Develop Credit Card Customer Churn Prediction. Journal of Theoretical and Applied Electronic Commerce Research. 2022; 17(4):1529-1542.
<https://doi.org/10.3390/jtaer17040077>
2. Siddiqui, N., Haque, M.A., Khan, S.M.S. et al. Different ML-based strategies for customer churn prediction in banking sector. J. of Data, Inf. and Manag. 6, 217–234 (2024).
<https://doi.org/10.1007/s42488-024-00126-z>
3. Peng, H., Zhang, L., & Liu, Q. (2023). Handling data imbalance in customer churn prediction for credit card users: An XGBoost approach. PLOS ONE, 18(8), e0289724.
<https://doi.org/10.1371/journal.pone.0289724>

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix,
roc_auc_score
import matplotlib.pyplot as plt
```

```
# Load the dataset
```

```
file_path = 'BankChurners.csv'
data = pd.read_csv(file_path)
```

```
data_info = data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10127 entries, 0 to 10126
```

```
Data columns (total 23 columns):
```

```
#   Column
Non-Null Count  Dtype
---  -
0   CLIENTNUM
10127 non-null  int64
1   Attrition_Flag
10127 non-null  object
2   Customer_Age
10127 non-null  int64
3   Gender
10127 non-null  object
4   Dependent_count
10127 non-null  int64
5   Education_Level
10127 non-null  object
6   Marital_Status
10127 non-null  object
7   Income_Category
10127 non-null  object
8   Card_Category
10127 non-null  object
9   Months_on_book
10127 non-null  int64
10  Total_Relationship_Count
10127 non-null  int64
11  Months_Inactive_12_mon
10127 non-null  int64
12  Contacts_Count_12_mon
10127 non-null  int64
13  Credit_Limit
10127 non-null  float64
14  Total_Revolving_Bal
```

```

10127 non-null int64
15 Avg_Open_To_Buy
10127 non-null float64
16 Total_Amt_Chng_Q4_Q1
10127 non-null float64
17 Total_Trans_Amt
10127 non-null int64
18 Total_Trans_Ct
10127 non-null int64
19 Total_Ct_Chng_Q4_Q1
10127 non-null float64
20 Avg_Utilization_Ratio
10127 non-null float64
21
Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1 10127 non-null float64
22
Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2 10127 non-null float64
dtypes: float64(7), int64(10), object(6)
memory usage: 1.8+ MB

data.drop(['CLIENTNUM',

'Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1',

'Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2'],
         axis=1,inplace=True)

missing_values = data.isnull().sum()
missing_values.sum()

0

numerical_summary = data.describe()
numerical_summary

```

	Customer_Age	Dependent_count	Months_on_book	\
count	10127.000000	10127.000000	10127.000000	
mean	46.325960	2.346203	35.928409	
std	8.016814	1.298908	7.986416	
min	26.000000	0.000000	13.000000	
25%	41.000000	1.000000	31.000000	
50%	46.000000	2.000000	36.000000	
75%	52.000000	3.000000	40.000000	
max	73.000000	5.000000	56.000000	

	Total_Relationship_Count	Months_Inactive_12_mon \
count	10127.000000	10127.000000
mean	3.812580	2.341167
std	1.554408	1.010622
min	1.000000	0.000000
25%	3.000000	2.000000
50%	4.000000	2.000000
75%	5.000000	3.000000
max	6.000000	6.000000

	Contacts_Count_12_mon	Credit_Limit	Total_Revolving_Bal \
count	10127.000000	10127.000000	10127.000000
mean	2.455317	8631.953698	1162.814061
std	1.106225	9088.776650	814.987335
min	0.000000	1438.300000	0.000000
25%	2.000000	2555.000000	359.000000
50%	2.000000	4549.000000	1276.000000
75%	3.000000	11067.500000	1784.000000
max	6.000000	34516.000000	2517.000000

	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt
Total_Trans_Ct \			
count	10127.000000	10127.000000	10127.000000
10127.000000			
mean	7469.139637	0.759941	4404.086304
64.858695			
std	9090.685324	0.219207	3397.129254
23.472570			
min	3.000000	0.000000	510.000000
10.000000			
25%	1324.500000	0.631000	2155.500000
45.000000			
50%	3474.000000	0.736000	3899.000000
67.000000			
75%	9859.000000	0.859000	4741.000000
81.000000			
max	34516.000000	3.397000	18484.000000
139.000000			

	Total_Ct_Chng_Q4_Q1	Avg_Utilization_Ratio
count	10127.000000	10127.000000
mean	0.712222	0.274894
std	0.238086	0.275691
min	0.000000	0.000000
25%	0.582000	0.023000
50%	0.702000	0.176000
75%	0.818000	0.503000
max	3.714000	0.999000

Numerical Summary

a. Wide Ranges

- Features like Credit_Limit and Total_Trans_Amt have large maximum values relative to the mean, suggesting the presence of outliers.

b. High Variability


- Columns like Total_Amt_Chng_Q4_Q1 (mean = 0.76, max = 3.39) and Total_Ct_Chng_Q4_Q1 (mean = 0.71, max = 3.71) indicate some extreme customer behaviors.

```
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Bar Plot 1: Age Distribution
data['Customer_Age'].hist(bins=15, ax=axes[0], color='orange',
edgecolor='black')
axes[0].set_title('Age Distribution', fontsize=14)
axes[0].set_xlabel('Age', fontsize=12)
axes[0].set_ylabel('Frequency', fontsize=12)

# Bar Plot 2: Gender Distribution
data['Gender'].value_counts().plot(kind='bar', ax=axes[1],
color=['#ba8087', '#74b6fc'], edgecolor='black')
axes[1].set_title('Gender Distribution', fontsize=14)
axes[1].set_xlabel('Gender', fontsize=12)
axes[1].set_ylabel('Count', fontsize=12)
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=0)

plt.tight_layout()
plt.show()
```



```
education_order = ['Uneducated', 'High School', 'College', 'Graduate',
'Post-Graduate', 'Doctorate', 'Unknown']
```

```


fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Bar Plot 1: Education Level Distribution (Original Order)
data['Education_Level'].value_counts()
[education_order].plot(kind='bar', ax=axes[0], color='#87CEEB',
edgecolor='black')
axes[0].set_title('Education Level Distribution', fontsize=14)
axes[0].set_xlabel('Education Level', fontsize=12)
axes[0].set_ylabel('Count', fontsize=12)
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=45,
fontsize=10)

# Bar Plot 2: Income Category Distribution
data['Income_Category'].value_counts().plot(kind='bar', ax=axes[1],
color='#FFA07A', edgecolor='black')
axes[1].set_title('Income Category Distribution', fontsize=14)
axes[1].set_xlabel('Income Category', fontsize=12)
axes[1].set_ylabel('Count', fontsize=12)
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=45,
fontsize=10)

plt.tight_layout()
plt.show()

```




```

plt.figure(figsize=(5, 5))
colors = ['#CE6A85', '#D6E3F8', '#FEF5EF', '#E4BB97']
data['Marital_Status'].value_counts().plot(kind='pie', autopct='%1.1f%%',
colors=colors, startangle=90)
plt.title('Marital Status Distribution', fontsize=14)
plt.ylabel('')
plt.show()

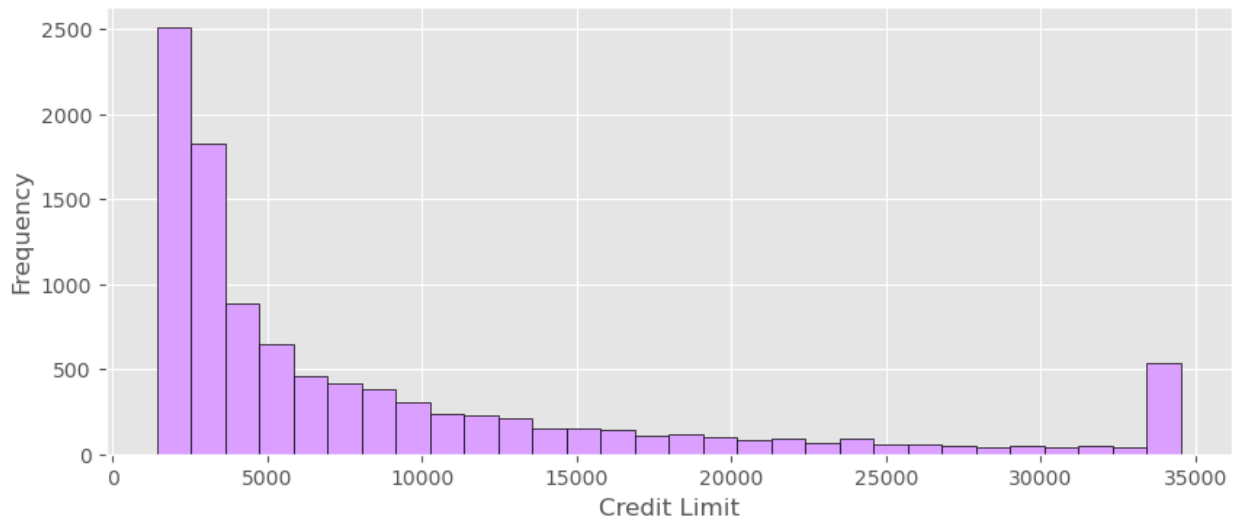
```

Marital Status Distribution



```
plt.figure(figsize=(10, 4))
data['Credit_Limit'].hist(bins=30, color='#db9fff', edgecolor='black')
plt.title('Distribution of Credit Limit', fontsize=14)
plt.xlabel('Credit Limit', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.show()
```

Distribution of Credit Limit




Summary

The dataset represents a middle-aged demographic with slightly more females than males, predominantly graduates, and customers in lower-to-middle income brackets. Most customers are married or single, and credit limits tend to be on the lower end.

```
# Pie Chart for Attrition Flag Distribution with numbers displayed
plt.figure(figsize=(5, 4))
attrition_counts = data['Attrition_Flag'].value_counts()
colors = ['#68c22d', '#a82246']

# Plotting the pie chart
plt.pie(attrition_counts, labels=attrition_counts.index,
autopct=lambda p: f'{p:.1f}%\n({int(p * attrition_counts.sum() / 100)})',
        startangle=90, colors=colors, textprops={'fontsize': 12})

plt.title('Attrition Flag Distribution', fontsize=14)
plt.show()
```



1. Proportion of Existing vs. Attrited Customers:
 - The chart shows that 83.9% (8,500) of the customers are existing customers, while 16.1% (1,627) are attrited customers.
 - This indicates a significant class imbalance in the dataset, with the majority of customers being retained.
1. Focus for Analysis:
 - The attrited customers represent a potential loss of revenue, and understanding their reasons for leaving could help mitigate future attrition.

- Efforts should focus on profiling and understanding the behaviors and characteristics of the attrited customers to identify:
- Key factors driving attrition (e.g., low credit limits, high utilization ratios, low transaction counts). Potential strategies for early intervention and retention.

```
numeric_columns = data.select_dtypes(include=['number']).columns

grouped_means = data.groupby('Attrition_Flag')[numeric_columns].mean()
grouped_means = pd.DataFrame(grouped_means)
```

```
from IPython.display import display
display(grouped_means.round(3))
```

	Customer_Age	Dependent_count	Months_on_book	\
Attrition_Flag				
Attrited Customer	46.659	2.403	36.178	
Existing Customer	46.262	2.335	35.881	

	Total_Relationship_Count	Months_Inactive_12_mon	\
Attrition_Flag			
Attrited Customer	3.280	2.693	
Existing Customer	3.915	2.274	

	Contacts_Count_12_mon	Credit_Limit
Total_Revolving_Bal		
Attrition_Flag		
Attrited Customer	2.972	8136.039
672.823		
Existing Customer	2.356	8726.878
1256.604		

	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1
Total_Trans_Amt		
Attrition_Flag		
Attrited Customer	7463.216	0.694
3095.026		
Existing Customer	7470.273	0.773
4654.656		

	Total_Trans_Ct	Total_Ct_Chng_Q4_Q1
Avg_Utilization_Ratio		
Attrition_Flag		
Attrited Customer	44.934	0.554
0.162		
Existing Customer	68.673	0.742
0.296		

```

def plot_attrition_proportion_for_numeric(column, bins, labels, title,
color):
    data[f'{column}_Group'] = pd.cut(data[column], bins=bins,
labels=labels)
    attrition = data.groupby([f'{column}_Group',
'Attrition_Flag']).size().unstack(fill_value=0)
    # Calculate proportions
    attrition_proportion = attrition.div(attrition.sum(axis=1),
axis=0)

    # Plot the proportion of attrited customers
    attrition_proportion['Attrited Customer'].plot(kind='bar',
color=color, edgecolor='black')
    plt.title(f'Proportion of Attrited Customers by {title}',
fontsize=14)
    plt.xlabel(title, fontsize=12)
    plt.ylabel('Proportion Attrited', fontsize=12)
    plt.ylim(0, 1)
    plt.xticks(rotation=45, fontsize=10)

# Setting up the plots
fig, axes = plt.subplots(3, 2, figsize=(18, 18))
plt.subplots_adjust(hspace=0.4, wspace=0.4)

# Total_Revolving_Bal
plt.sca(axes[0, 0])
plot_attrition_proportion_for_numeric('Total_Revolving_Bal',
bins=[0, 500, 1000, 1500, 2000,
2500, 3000],
labels=['<500', '500-1000',
'1000-1500', '1500-2000', '2000-2500', '>2500'],
title='Total Revolving Balance',
color='#FFA07A')

# Total_Trans_Amt
plt.sca(axes[0, 1])
plot_attrition_proportion_for_numeric('Total_Trans_Amt',
bins=[0, 2000, 4000, 6000, 8000,
10000, 20000],
labels=['<2K', '2K-4K', '4K-6K',
'6K-8K', '8K-10K', '>10K'],
title='Total Transaction
Amount',
color='#87CEEB')

# Total_Trans_Ct
plt.sca(axes[1, 0])
plot_attrition_proportion_for_numeric('Total_Trans_Ct',
bins=[0, 20, 40, 60, 80, 100,

```

```

140],
'60-80', '80-100', '>100'],
                                labels=['<20', '20-40', '40-60',
                                title='Total Transaction Count',
                                color='#FFD700')

# Total_Ct_Chng_Q4_Q1
plt.sca(axes[1, 1])
plot_attrition_proportion_for_numeric('Total_Ct_Chng_Q4_Q1',
bins=[0, 0.5, 1.0, 1.5, 2.0,
2.5, 3.5],
                                labels=['<0.5', '0.5-1.0', '1.0-
1.5', '1.5-2.0', '2.0-2.5', '>2.5'],
                                title='Total Count Change Q4-
Q1',
                                color='#32CD32')

# Avg_Utilization_Ratio
plt.sca(axes[2, 0])
plot_attrition_proportion_for_numeric('Avg_Utilization_Ratio',
bins=[0, 0.2, 0.4, 0.6, 0.8,
1.0],
                                labels=['0-0.2', '0.2-0.4',
'0.4-0.6', '0.6-0.8', '0.8-1.0'],
                                title='Average Utilization
Ratio',
                                color='#FF6347')

# Hiding the last empty subplot
axes[2, 1].axis('off')

plt.show()

```

/tmp/ipykernel_4745/4239478096.py:7: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

attrition = data.groupby([f'{column}_Group',
'Attrition_Flag']).size().unstack(fill_value=0)

```

/tmp/ipykernel_4745/4239478096.py:7: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.






```

attrition = data.groupby([f'{column}_Group',
'Attrition_Flag']).size().unstack(fill_value=0)

```

/tmp/ipykernel_4745/4239478096.py:7: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
attrition = data.groupby([f'{column}_Group',
'Attrition_Flag']).size().unstack(fill_value=0)
/tmp/ipykernel_4745/4239478096.py:7: FutureWarning: The default of
observed=False is deprecated and will be changed to True in a future
version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
attrition = data.groupby([f'{column}_Group',
'Attrition_Flag']).size().unstack(fill_value=0)
/tmp/ipykernel_4745/4239478096.py:7: FutureWarning: The default of
observed=False is deprecated and will be changed to True in a future
version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
attrition = data.groupby([f'{column}_Group',
'Attrition_Flag']).size().unstack(fill_value=0)
```

Proportion Summary

- **Low Engagement Indicators:** Low transaction amounts, fewer transactions, and declines in transaction activity strongly correlate with attrition.
- **High Utilization Risk:** Customers with high utilization ratios are at risk of attrition, potentially due to financial strain.
- **Retention Focus:** Customers with low revolving balances or low transaction activity should be flagged for retention strategies.

```

import pandas as pd
from sklearn.preprocessing import MinMaxScaler

file_path = '/home/jovyan/Web_Recommender/BankChurners.csv'
data = pd.read_csv(file_path)

columns_to_drop = [
    "CLIENTNUM",

    "Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1",

    "Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2"
]
data_cleaned = data.drop(columns=columns_to_drop, errors='ignore')

data_cleaned.columns = data_cleaned.columns.str.strip()

categorical_columns = [
    "Attrition_Flag", "Gender", "Education_Level",
    "Marital_Status", "Income_Category", "Card_Category"
]
data_encoded = pd.get_dummies(data_cleaned,
columns=categorical_columns, drop_first=True)

numerical_columns = [
    "Customer_Age", "Dependent_count", "Months_on_book",
    "Total_Relationship_Count",
    "Months_Inactive_12_mon", "Contacts_Count_12_mon", "Credit_Limit",

    "Total_Revolving_Bal", "Avg_Open_To_Buy", "Total_Amt_Chng_Q4_Q1",
    "Total_Trans_Amt", "Total_Trans_Ct", "Total_Ct_Chng_Q4_Q1",
    "Avg_Utilization_Ratio"
]
scaler = MinMaxScaler()
data_encoded[numerical_columns] =
scaler.fit_transform(data_encoded[numerical_columns])

for col in numerical_columns:
    Q1 = data_encoded[col].quantile(0.25)
    Q3 = data_encoded[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    data_encoded[col] = data_encoded[col].clip(lower=lower_bound,
upper=upper_bound)

data_encoded["Transactions_Per_Month"] = (
    data_encoded["Total_Trans_Amt"] / (data_encoded["Months_on_book"]

```

```

+ 1)
)

output_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv'
data_encoded.to_csv(output_path, index=False)

output_path
'/home/jovyan/Web_Recommender/Cleaned_dataset.csv'

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.preprocessing import StandardScaler

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Change this to
your target column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

logistic_model = LogisticRegression(random_state=42, max_iter=1000)

logistic_model.fit(X_train_scaled, y_train)

y_pred = logistic_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

coefficients = pd.DataFrame({
    'Feature': features.columns,
    'Coefficient': logistic_model.coef_[0]
})

```

```
).sort_values(by='Coefficient', ascending=False)
```

```
print("\nLogistic Regression Coefficients:")
```

```
print(coefficients)
```

Classification Report:

	precision	recall	f1-score	support
False	0.78	0.62	0.69	325
True	0.93	0.97	0.95	1701
accuracy			0.91	2026
macro avg	0.85	0.79	0.82	2026
weighted avg	0.91	0.91	0.91	2026

Accuracy Score: 0.9106614017769002

Confusion Matrix:

```
[[ 201  124]
 [   57 1644]]
```

Logistic Regression Coefficients:

	Feature	Coefficient
11	Total_Trans_Ct	3.512005
12	Total_Ct_Chng_Q4_Q1	0.730180
7	Total_Revolving_Bal	0.708653
3	Total_Relationship_Count	0.601870
14	Gender_M	0.554394
27	Income_Category_Less than \$40K	0.502909
24	Income_Category_\$40K - \$60K	0.427671
28	Income_Category_Unknown	0.351166
21	Marital_Status_Married	0.291811
25	Income_Category_\$60K - \$80K	0.223433
8	Avg_Open_To_Buy	0.183869
26	Income_Category_\$80K - \$120K	0.182795
9	Total_Amt_Chng_Q4_Q1	0.141366
0	Customer_Age	0.095678
13	Avg_Utilization_Ratio	0.086499
23	Marital_Status_Unknown	0.046666
6	Credit_Limit	0.035938
22	Marital_Status_Single	0.019190
16	Education_Level_Graduate	-0.007410
17	Education_Level_High School	-0.032070
30	Card_Category_Platinum	-0.034264
29	Card_Category_Gold	-0.034658
2	Months_on_book	-0.037354
19	Education_Level_Uneducated	-0.047129
20	Education_Level_Unknown	-0.053187
18	Education_Level_Post-Graduate	-0.061149
31	Card_Category_Silver	-0.076589
15	Education_Level_Doctorate	-0.093723

1	Dependent_count	-0.155334
32	Transactions_Per_Month	-0.398914
5	Contacts_Count_12_mon	-0.580440
4	Months_Inactive_12_mon	-0.599302
10	Total_Trans_Amt	-1.790278

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.preprocessing import StandardScaler

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Change this to
your target column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn_model = KNeighborsClassifier(n_neighbors=5) # You can experiment
with the number of neighbors (n_neighbors)

knn_model.fit(X_train_scaled, y_train)

y_pred = knn_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
False	0.73	0.40	0.52	325
True	0.89	0.97	0.93	1701

accuracy			0.88	2026
macro avg	0.81	0.69	0.73	2026
weighted avg	0.87	0.88	0.87	2026

Accuracy Score: 0.8800592300098716

Confusion Matrix:

```
[[ 131  194]
 [   49 1652]]
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.tree import export_text

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Corrected target
column name
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

dt_model = DecisionTreeClassifier(random_state=42, max_depth=5) #
Adjust max_depth for complexity

dt_model.fit(X_train, y_train)

y_pred = dt_model.predict(X_test)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

tree_rules = export_text(dt_model,
feature_names=list(features.columns))
print("\nDecision Tree Rules:")
print(tree_rules)
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

False	0.80	0.81	0.80	325
True	0.96	0.96	0.96	1701
accuracy			0.94	2026
macro avg	0.88	0.89	0.88	2026
weighted avg	0.94	0.94	0.94	2026

Accuracy Score: 0.9363277393879565

Confusion Matrix:

```
[[ 264  61]
 [  68 1633]]
```

Decision Tree Rules:

```
|--- Total_Trans_Ct <= 0.34
|   |--- Total_Revolving_Bal <= 0.24
|   |   |--- Total_Ct_Chng_Q4_Q1 <= 0.18
|   |   |   |--- Months_Inactive_12_mon <= 0.25
|   |   |   |   |--- Total_Trans_Amt <= 0.09
|   |   |   |   |   |--- class: True
|   |   |   |   |--- Total_Trans_Amt > 0.09
|   |   |   |   |   |--- class: False
|   |   |   |--- Months_Inactive_12_mon > 0.25
|   |   |   |   |--- Total_Trans_Amt <= 0.07
|   |   |   |   |   |--- class: False
|   |   |   |   |--- Total_Trans_Amt > 0.07
|   |   |   |   |   |--- class: False
|   |   |--- Total_Ct_Chng_Q4_Q1 > 0.18
|   |   |   |--- Total_Relationship_Count <= 0.30
|   |   |   |   |--- Card_Category_Gold <= 0.50
|   |   |   |   |   |--- class: False
|   |   |   |   |--- Card_Category_Gold > 0.50
|   |   |   |   |   |--- class: True
|   |   |   |--- Total_Relationship_Count > 0.30
|   |   |   |   |--- Total_Trans_Amt <= 0.08
|   |   |   |   |   |--- class: True
|   |   |   |   |--- Total_Trans_Amt > 0.08
|   |   |   |   |   |--- class: False
|   |--- Total_Revolving_Bal > 0.24
|   |   |--- Total_Relationship_Count <= 0.30
|   |   |   |--- Total_Ct_Chng_Q4_Q1 <= 0.26
|   |   |   |   |--- Contacts_Count_12_mon <= 0.13
|   |   |   |   |   |--- class: True
|   |   |   |   |--- Contacts_Count_12_mon > 0.13
|   |   |   |   |   |--- class: False
|   |   |   |--- Total_Ct_Chng_Q4_Q1 > 0.26
|   |   |   |   |--- Total_Trans_Amt <= 0.09
|   |   |   |   |   |--- class: True
|   |   |   |   |--- Total_Trans_Amt > 0.09
|   |   |   |   |   |--- class: False
```

```

--- Total_Relationship_Count > 0.30
--- Total_Trans_Amt <= 0.09
|--- Transactions_Per_Month <= 0.02
|   |--- class: False
|--- Transactions_Per_Month > 0.02
|   |--- class: True
--- Total_Trans_Amt > 0.09
|--- Total_Ct_Chng_Q4_Q1 <= 0.17
|   |--- class: False
|--- Total_Ct_Chng_Q4_Q1 > 0.17
|   |--- class: True
--- Total_Trans_Ct > 0.34
|--- Total_Trans_Amt <= 0.27
|   |--- Total_Trans_Ct <= 0.37
|       |--- Total_Relationship_Count <= 0.30
|       |--- Total_Ct_Chng_Q4_Q1 <= 0.24
|       |   |--- class: False
|       |--- Total_Ct_Chng_Q4_Q1 > 0.24
|       |   |--- class: False
|   --- Total_Relationship_Count > 0.30
|       |--- Total_Trans_Amt <= 0.23
|       |   |--- class: True
|       |--- Total_Trans_Amt > 0.23
|       |   |--- class: False
|   --- Total_Trans_Ct > 0.37
|       |--- Total_Trans_Ct <= 0.41
|           |--- Total_Revolving_Bal <= 0.16
|           |   |--- class: True
|           |--- Total_Revolving_Bal > 0.16
|           |   |--- class: True
|       --- Total_Trans_Ct > 0.41
|           |--- Card_Category_Platinum <= 0.50
|           |   |--- class: True
|           |--- Card_Category_Platinum > 0.50
|           |   |--- class: False
|   --- Total_Trans_Amt > 0.27
|       |--- Total_Trans_Ct <= 0.53
|           |--- Total_Revolving_Bal <= 0.34
|           |   |--- Credit_Limit <= 0.04
|           |   |   |--- class: True
|           |   |--- Credit_Limit > 0.04
|           |   |   |--- class: False
|       --- Total_Revolving_Bal > 0.34
|           |--- Total_Amt_Chng_Q4_Q1 <= 0.26
|           |   |--- class: True
|           |--- Total_Amt_Chng_Q4_Q1 > 0.26
|           |   |--- class: False
|   --- Total_Trans_Ct > 0.53
|       |--- Total_Trans_Ct <= 0.60

```



```
| | | | | --- Total_Trans_Amt <= 0.44
| | | | | | --- class: True
| | | | | --- Total_Trans_Amt > 0.44
| | | | | | --- class: False
| | | | | --- Total_Trans_Ct > 0.60
| | | | | | --- Total_Ct_Chng_Q4_Q1 <= 0.25
| | | | | | --- class: True
| | | | | --- Total_Ct_Chng_Q4_Q1 > 0.25
| | | | | | --- class: True
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update this with the correct path to your preprocessed dataset
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Change this to
your target column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

rf_classifier = RandomForestClassifier(n_estimators=100,
random_state=42)

rf_classifier.fit(X_train, y_train)

y_pred = rf_classifier.predict(X_test)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

feature_importances = pd.DataFrame({
    'Feature': features.columns,
    'Importance': rf_classifier.feature_importances_
}).sort_values(by='Importance', ascending=False)
```

```
print("\nFeature Importances:")
print(feature_importances)
```

Classification Report:

	precision	recall	f1-score	support
False	0.93	0.82	0.87	325
True	0.97	0.99	0.98	1701
accuracy			0.96	2026
macro avg	0.95	0.90	0.92	2026
weighted avg	0.96	0.96	0.96	2026

Accuracy Score: 0.9610069101678184

Confusion Matrix:

```
[[ 265  60]
 [  19 1682]]
```

Feature Importances:

	Feature	Importance
11	Total_Trans_Ct	0.142320
10	Total_Trans_Amt	0.128612
32	Transactions_Per_Month	0.122163
12	Total_Ct_Chng_Q4_Q1	0.106575
7	Total_Revolving_Bal	0.102195
9	Total_Amt_Chng_Q4_Q1	0.056982
13	Avg_Utilization_Ratio	0.056403
3	Total_Relationship_Count	0.056097
0	Customer_Age	0.033364
6	Credit_Limit	0.029344
8	Avg_Open_To_Buy	0.027856
2	Months_on_book	0.024362
4	Months_Inactive_12_mon	0.022898
5	Contacts_Count_12_mon	0.021790
1	Dependent_count	0.013752
14	Gender_M	0.008632
21	Marital_Status_Married	0.005864
22	Marital_Status_Single	0.004623
27	Income_Category_Less than \$40K	0.003727
16	Education_Level_Graduate	0.003555
17	Education_Level_High School	0.003200
20	Education_Level_Unknown	0.003151
25	Income_Category_\$60K - \$80K	0.002827
19	Education_Level_Uneducated	0.002805
26	Income_Category_\$80K - \$120K	0.002779
24	Income_Category_\$40K - \$60K	0.002635
23	Marital_Status_Unknown	0.002301
28	Income_Category_Unknown	0.002103
18	Education_Level_Post-Graduate	0.002081
15	Education_Level_Doctorate	0.001929

31	Card_Category_Silver	0.001666
29	Card_Category_Gold	0.001164
30	Card_Category_Platinum	0.000245

```
!pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in
```

```
/opt/conda/lib/python3.11/site-packages (0.12.4)
```

```
Requirement already satisfied: numpy>=1.17.3 in
```

```
/opt/conda/lib/python3.11/site-packages (from imbalanced-learn)
(1.26.4)
```

```
Requirement already satisfied: scipy>=1.5.0 in
```

```
/opt/conda/lib/python3.11/site-packages (from imbalanced-learn)
(1.12.0)
```

```
Requirement already satisfied: scikit-learn>=1.0.2 in
```

```
/opt/conda/lib/python3.11/site-packages (from imbalanced-learn)
(1.5.1)
```

```
Requirement already satisfied: joblib>=1.1.1 in
```

```
/opt/conda/lib/python3.11/site-packages (from imbalanced-learn)
(1.4.2)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in
```

```
/opt/conda/lib/python3.11/site-packages (from imbalanced-learn)
(3.5.0)
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
```

```
from imblearn.over_sampling import SMOTE
```

```
from sklearn.preprocessing import StandardScaler
```

```
file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
```

```
Update with your dataset path
```

```
data = pd.read_csv(file_path)
```

```
target_column = 'Attrition_Flag_Existing Customer' # Update this with
your target column name
```

```
features = data.drop(columns=[target_column]) # Remove the target
column from features
```

```
target = data[target_column]
```

```
X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
smote = SMOTE(random_state=42)
```

```

X_train_smote, y_train_smote = smote.fit_resample(X_train_scaled,
y_train)

rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

rf_model.fit(X_train_smote, y_train_smote)

y_pred = rf_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

feature_importances = pd.DataFrame({
    'Feature': features.columns,
    'Importance': rf_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importances:")
print(feature_importances)

```

```

Classification Report:

```

	precision	recall	f1-score	support
False	0.87	0.90	0.89	325
True	0.98	0.97	0.98	1701
accuracy			0.96	2026
macro avg	0.93	0.94	0.93	2026
weighted avg	0.96	0.96	0.96	2026

```

Accuracy Score: 0.9629812438302073

```

```

Confusion Matrix:

```

```

[[ 293  32]
 [  43 1658]]

```

```

Feature Importances:

```

	Feature	Importance
11	Total_Trans_Ct	0.182145
10	Total_Trans_Amt	0.143929
32	Transactions_Per_Month	0.128141
7	Total_Revolving_Bal	0.085220
12	Total_Ct_Chng_Q4_Q1	0.075869
3	Total_Relationship_Count	0.054354
4	Months_Inactive_12_mon	0.050511
9	Total_Amt_Chng_Q4_Q1	0.045538
13	Avg_Utilization_Ratio	0.044740

5	Contacts_Count_12_mon	0.042764
0	Customer_Age	0.024315
6	Credit_Limit	0.019517
2	Months_on_book	0.018512
8	Avg_Open_To_Buy	0.018491
1	Dependent_count	0.017042
14	Gender_M	0.010819
21	Marital_Status_Married	0.005848
22	Marital_Status_Single	0.005107
16	Education_Level_Graduate	0.003113
27	Income_Category_Less than \$40K	0.002932
26	Income_Category_\$80K - \$120K	0.002914
19	Education_Level_Uneducated	0.002333
17	Education_Level_High School	0.002300
24	Income_Category_\$40K - \$60K	0.002207
25	Income_Category_\$60K - \$80K	0.001956
20	Education_Level_Unknown	0.001893
18	Education_Level_Post-Graduate	0.001628
28	Income_Category_Unknown	0.001492
23	Marital_Status_Unknown	0.001309
31	Card_Category_Silver	0.001116
15	Education_Level_Doctorate	0.001011
29	Card_Category_Gold	0.000805
30	Card_Category_Platinum	0.000128

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.preprocessing import StandardScaler

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Update this with
your target column name
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

gb_model = GradientBoostingClassifier(n_estimators=100,
```

```

learning_rate=0.1, max_depth=3, random_state=42)

gb_model.fit(X_train_scaled, y_train)

y_pred = gb_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

feature_importances = pd.DataFrame({
    'Feature': features.columns,
    'Importance': gb_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importances:")
print(feature_importances)

```

Classification Report:

	precision	recall	f1-score	support
False	0.93	0.85	0.89	325
True	0.97	0.99	0.98	1701
accuracy			0.97	2026
macro avg	0.95	0.92	0.93	2026
weighted avg	0.96	0.97	0.96	2026

Accuracy Score: 0.9654491609081934

Confusion Matrix:

```

[[ 276  49]
 [  21 1680]]

```

Feature Importances:

	Feature	Importance
11	Total_Trans_Ct	0.333336
10	Total_Trans_Amt	0.186573
7	Total_Revolving_Bal	0.166748
12	Total_Ct_Chng_Q4_Q1	0.118200
3	Total_Relationship_Count	0.094148
9	Total_Amt_Chng_Q4_Q1	0.033666
32	Transactions_Per_Month	0.018806
4	Months_Inactive_12_mon	0.015172
0	Customer_Age	0.014258
5	Contacts_Count_12_mon	0.007577
8	Avg_Open_To_Buy	0.002917
2	Months_on_book	0.001867

14	Gender_M	0.001674
21	Marital_Status_Married	0.001650
1	Dependent_count	0.001225
13	Avg_Utilization_Ratio	0.000766
30	Card_Category_Platinum	0.000631
29	Card_Category_Gold	0.000475
6	Credit_Limit	0.000161
22	Marital_Status_Single	0.000095
16	Education_Level_Graduate	0.000054
20	Education_Level_Unknown	0.000002
19	Education_Level_Uneducated	0.000000
15	Education_Level_Doctorate	0.000000
24	Income_Category_\$40K - \$60K	0.000000
25	Income_Category_\$60K - \$80K	0.000000
26	Income_Category_\$80K - \$120K	0.000000
27	Income_Category_Less than \$40K	0.000000
28	Income_Category_Unknown	0.000000
18	Education_Level_Post-Graduate	0.000000
17	Education_Level_High School	0.000000
31	Card_Category_Silver	0.000000
23	Marital_Status_Unknown	0.000000

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.preprocessing import StandardScaler
from lightgbm import LGBMClassifier

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Correct target
column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

lgbm_model = LGBMClassifier(
    objective='binary',
    boosting_type='gbdt',
    learning_rate=0.1,
```

```

    n_estimators=100,
    num_leaves=31,
    random_state=42
)

lgbm_model.fit(
    X_train_scaled,
    y_train,
    eval_set=[(X_test_scaled, y_test)],
    eval_metric='logloss'
)

y_pred = lgbm_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

feature_importances = pd.DataFrame({
    'Feature': features.columns,
    'Importance': lgbm_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importances:")
print(feature_importances)

[LightGBM] [Info] Number of positive: 6799, number of negative: 1302
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.005755 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2322
[LightGBM] [Info] Number of data points in the train set: 8101, number
of used features: 32
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.839279 ->
initscore=1.652874
[LightGBM] [Info] Start training from score 1.652874
Classification Report:

```

	precision	recall	f1-score	support
False	0.91	0.91	0.91	325
True	0.98	0.98	0.98	1701
accuracy			0.97	2026
macro avg	0.95	0.95	0.95	2026
weighted avg	0.97	0.97	0.97	2026

Accuracy Score: 0.9718657453109576

Confusion Matrix:

```
[[ 297   28]
 [   29 1672]]
```

Feature Importances:

	Feature	Importance
10	Total_Trans_Amt	543
11	Total_Trans_Ct	362
9	Total_Amt_Chng_Q4_Q1	350
12	Total_Ct_Chng_Q4_Q1	238
32	Transactions_Per_Month	205
0	Customer_Age	179
7	Total_Revolving_Bal	172
3	Total_Relationship_Count	153
6	Credit_Limit	135
5	Contacts_Count_12_mon	110
2	Months_on_book	107
8	Avg_Open_To_Buy	103
4	Months_Inactive_12_mon	92
13	Avg_Utilization_Ratio	66
21	Marital_Status_Married	40
1	Dependent_count	39
14	Gender_M	25
25	Income_Category_\$60K - \$80K	16
18	Education_Level_Post-Graduate	10
22	Marital_Status_Single	9
24	Income_Category_\$40K - \$60K	8
16	Education_Level_Graduate	7
20	Education_Level_Unknown	7
19	Education_Level_Uneducated	6
15	Education_Level_Doctorate	4
27	Income_Category_Less than \$40K	4
23	Marital_Status_Unknown	3
26	Income_Category_\$80K - \$120K	3
17	Education_Level_High School	2
28	Income_Category_Unknown	2
29	Card_Category_Gold	0
30	Card_Category_Platinum	0
31	Card_Category_Silver	0

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
```

```
file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
```

```

Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Correct target
column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

xgb_model = XGBClassifier(
    objective='binary:logistic',
    learning_rate=0.1,
    n_estimators=100,
    max_depth=5,
    random_state=42,
    use_label_encoder=False, # To avoid a warning for newer versions
of XGBoost
    eval_metric='logloss' # Specify evaluation metric
)

xgb_model.fit(X_train_scaled, y_train)

y_pred = xgb_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

feature_importances = pd.DataFrame({
    'Feature': features.columns,
    'Importance': xgb_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importances:")
print(feature_importances)

```

```

Classification Report:

```

	precision	recall	f1-score	support
False	0.92	0.90	0.91	325

True	0.98	0.99	0.98	1701
accuracy			0.97	2026
macro avg	0.95	0.94	0.95	2026
weighted avg	0.97	0.97	0.97	2026

Accuracy Score: 0.9713721618953604

Confusion Matrix:

```
[[ 292  33]
 [ 25 1676]]
```

Feature Importances:

	Feature	Importance
11	Total_Trans_Ct	0.241002
7	Total_Revolving_Bal	0.151338
3	Total_Relationship_Count	0.103535
10	Total_Trans_Amt	0.059479
12	Total_Ct_Chng_Q4_Q1	0.049363
4	Months_Inactive_12_mon	0.043459
8	Avg_Open_To_Buy	0.033538
9	Total_Amt_Chng_Q4_Q1	0.030702
0	Customer_Age	0.029823
32	Transactions_Per_Month	0.027283
14	Gender_M	0.027039
5	Contacts_Count_12_mon	0.022109
6	Credit_Limit	0.018385
21	Marital_Status_Married	0.016800
2	Months_on_book	0.014726
1	Dependent_count	0.013676
23	Marital_Status_Unknown	0.012103
13	Avg_Utilization_Ratio	0.011663
15	Education_Level_Doctorate	0.011065
28	Income_Category_Unknown	0.010864
16	Education_Level_Graduate	0.010569
25	Income_Category_\$60K - \$80K	0.010164
18	Education_Level_Post-Graduate	0.009461
19	Education_Level_Uneducated	0.007283
27	Income_Category_Less than \$40K	0.006160
22	Marital_Status_Single	0.005295
17	Education_Level_High School	0.005217
24	Income_Category_\$40K - \$60K	0.005143
26	Income_Category_\$80K - \$120K	0.004528
29	Card_Category_Gold	0.004250
20	Education_Level_Unknown	0.003978
30	Card_Category_Platinum	0.000000
31	Card_Category_Silver	0.000000

```
import pandas as pd
from sklearn.model_selection import train_test_split
```

```

from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Correct target
column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

xgb_model = XGBClassifier(
    objective='binary:logistic',
    learning_rate=0.1,
    n_estimators=100,
    max_depth=5,
    random_state=42,
    use_label_encoder=False,
    eval_metric='logloss'
)

rf_model = RandomForestClassifier(
    n_estimators=100,
    max_depth=5,
    random_state=42
)

lr_model = LogisticRegression(
    random_state=42,
    max_iter=1000
)

ensemble_model = VotingClassifier(
    estimators=[
        ('xgb', xgb_model),
        ('rf', rf_model),
        ('lr', lr_model)
    ]
)

```

```

    ],
    voting='hard' # Use 'soft' for probability-based voting
)

ensemble_model.fit(X_train_scaled, y_train)

y_pred = ensemble_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Classification Report:

	precision	recall	f1-score	support
False	0.92	0.71	0.80	325
True	0.95	0.99	0.97	1701
accuracy			0.94	2026
macro avg	0.93	0.85	0.88	2026
weighted avg	0.94	0.94	0.94	2026

Accuracy Score: 0.9432379072063178

Confusion Matrix:

```

[[ 230   95]
 [   20 1681]]

```

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
from xgboost import XGBClassifier
import numpy as np

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Correct target
column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

features['Credit_to_Income_Ratio'] = features['Credit_Limit'] /
(features['Avg_Open_To_Buy'] + 1)

```

```

features['Transaction_Rate'] = features['Total_Trans_Amt'] /
(features['Months_on_book'] + 1)

features['High_Utilization'] = (features['Avg_Utilization_Ratio'] >
0.5).astype(int)

for col in ['Total_Trans_Amt', 'Avg_Open_To_Buy']:
    features[f'Log_{col}'] = np.log1p(features[col])

categorical_cols = features.select_dtypes(include=['object']).columns
if len(categorical_cols) > 0:
    encoder = OneHotEncoder(sparse=False, drop='first')
    encoded_cols = encoder.fit_transform(features[categorical_cols])
    encoded_features = pd.DataFrame(encoded_cols,
columns=encoder.get_feature_names_out(categorical_cols))
    features = pd.concat([features.drop(columns=categorical_cols),
encoded_features], axis=1)

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

xgb_model = XGBClassifier(
    objective='binary:logistic',
    learning_rate=0.1,
    n_estimators=100,
    max_depth=5,
    random_state=42,
    use_label_encoder=False,
    eval_metric='logloss'
)

xgb_model.fit(X_train_scaled, y_train)

y_pred = xgb_model.predict(X_test_scaled)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

import matplotlib.pyplot as plt

feature_importances = pd.DataFrame({

```

```

    'Feature': features.columns,
    'Importance': xgb_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

```

```

print("\nFeature Importances:")
print(feature_importances)

```

```

plt.figure(figsize=(10, 6))
plt.barh(feature_importances['Feature'],
feature_importances['Importance'])
plt.xlabel('Feature Importance')
plt.ylabel('Feature')
plt.title('Feature Importances from XGBoost')
plt.gca().invert_yaxis()
plt.show()

```

Classification Report:

	precision	recall	f1-score	support
False	0.91	0.90	0.91	325
True	0.98	0.98	0.98	1701
accuracy			0.97	2026
macro avg	0.95	0.94	0.94	2026
weighted avg	0.97	0.97	0.97	2026

Accuracy Score: 0.9703849950641659

Confusion Matrix:

```

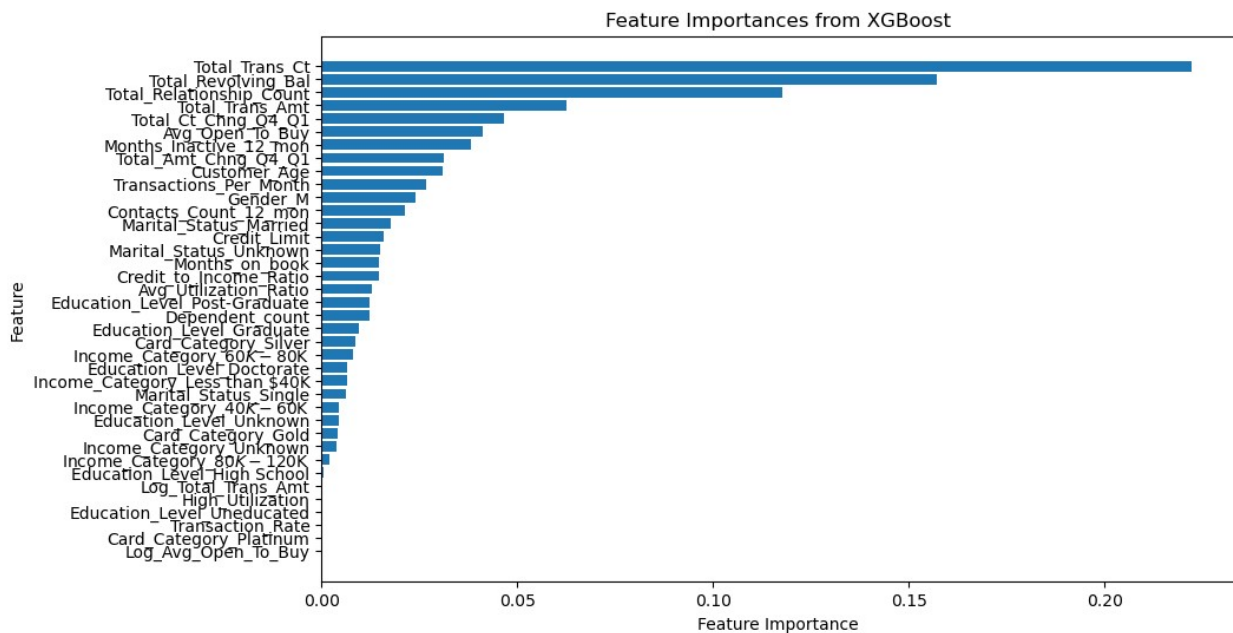
[[ 293  32]
 [ 28 1673]]

```

Feature Importances:

	Feature	Importance
11	Total_Trans_Ct	0.222327
7	Total_Revolving_Bal	0.157008
3	Total_Relationship_Count	0.117635
10	Total_Trans_Amt	0.062698
12	Total_Ct_Chng_Q4_Q1	0.046681
8	Avg_Open_To_Buy	0.041149
4	Months_Inactive_12_mon	0.038210
9	Total_Amt_Chng_Q4_Q1	0.031182
0	Customer_Age	0.031005
32	Transactions_Per_Month	0.026838
14	Gender_M	0.024096
5	Contacts_Count_12_mon	0.021386
21	Marital_Status_Married	0.017549
6	Credit_Limit	0.015906
23	Marital_Status_Unknown	0.014859
2	Months_on_book	0.014714
33	Credit_to_Income_Ratio	0.014592

13	Avg_Utilization_Ratio	0.012785
18	Education_Level_Post-Graduate	0.012311
1	Dependent_count	0.012289
16	Education_Level_Graduate	0.009521
31	Card_Category_Silver	0.008518
25	Income_Category_\$60K - \$80K	0.007960
15	Education_Level_Doctorate	0.006525
27	Income_Category_Less than \$40K	0.006424
22	Marital_Status_Single	0.006138
24	Income_Category_\$40K - \$60K	0.004566
20	Education_Level_Unknown	0.004320
29	Card_Category_Gold	0.004271
28	Income_Category_Unknown	0.003815
26	Income_Category_\$80K - \$120K	0.002059
17	Education_Level_High School	0.000665
36	Log_Total_Trans_Amt	0.000000
35	High_Utilization	0.000000
19	Education_Level_Uneducated	0.000000
34	Transaction_Rate	0.000000
30	Card_Category_Platinum	0.000000
37	Log_Avg_Open_To_Buy	0.000000



```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import RFE
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix
```



```

from xgboost import XGBClassifier

file_path = '/home/jovyan/Web_Recommender/Cleaned_dataset.csv' #
Update with your dataset path
data = pd.read_csv(file_path)

target_column = 'Attrition_Flag_Existing Customer' # Correct target
column
features = data.drop(columns=[target_column]) # Remove the target
column from features
target = data[target_column]

X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42, stratify=target)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

xgb_model = XGBClassifier(
    objective='binary:logistic',
    learning_rate=0.1,
    n_estimators=100,
    max_depth=5,
    random_state=42,
    use_label_encoder=False,
    eval_metric='logloss'
)

rfe = RFE(estimator=xgb_model, n_features_to_select=10, step=1) #
Select top 10 features
rfe.fit(X_train_scaled, y_train)

selected_features = features.columns[rfe.support_]
print("Selected Features:")
print(selected_features)

X_train_rfe = rfe.transform(X_train_scaled)
X_test_rfe = rfe.transform(X_test_scaled)

xgb_model.fit(X_train_rfe, y_train)

y_pred = xgb_model.predict(X_test_rfe)

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Accuracy Score:", accuracy_score(y_test, y_pred))

print("Confusion Matrix:")

```

```

print(confusion_matrix(y_test, y_pred))

import matplotlib.pyplot as plt

feature_importances = pd.DataFrame({
    'Feature': selected_features,
    'Importance': xgb_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importances for Selected Features:")
print(feature_importances)

plt.figure(figsize=(10, 6))
plt.barh(feature_importances['Feature'],
feature_importances['Importance'])
plt.xlabel('Feature Importance')
plt.ylabel('Feature')
plt.title('Feature Importances for Selected Features (RFE)')
plt.gca().invert_yaxis()
plt.show()

```

Selected Features:

```

Index(['Customer_Age', 'Total_Relationship_Count',
'Months_Inactive_12_mon',
'Contacts_Count_12_mon', 'Total_Revolving_Bal',
'Avg_Open_To_Buy',
'Total_Amt_Chng_Q4_Q1', 'Total_Trans_Amt', 'Total_Trans_Ct',
'Total_Ct_Chng_Q4_Q1'],
dtype='object')

```

Classification Report:

	precision	recall	f1-score	support
False	0.92	0.91	0.91	325
True	0.98	0.98	0.98	1701
accuracy			0.97	2026
macro avg	0.95	0.95	0.95	2026
weighted avg	0.97	0.97	0.97	2026

Accuracy Score: 0.9718657453109576

Confusion Matrix:

```

[[ 295   30]
 [   27 1674]]

```

Feature Importances for Selected Features:

	Feature	Importance
8	Total_Trans_Ct	0.293107
4	Total_Revolving_Bal	0.199824
1	Total_Relationship_Count	0.154497
7	Total_Trans_Amt	0.083097

9	Total_Ct_Chng_Q4_Q1	0.068072
2	Months_Inactive_12_mon	0.053755
6	Total_Amt_Chng_Q4_Q1	0.042817
0	Customer_Age	0.039844
5	Avg_Open_To_Buy	0.034047
3	Contacts_Count_12_mon	0.030942

